

TODO: Paper Title

Anonymous author

Anonymous affiliation

Abstract

Background. TODO: 1–2 sentences on the problem and why it matters.

Aims. TODO: the research questions or specific objectives.

Method. TODO: dataset, models, approaches compared.

Results. TODO: headline numbers, ordered by RQ.

Conclusions. TODO: what this means and the takeaway for practitioners.

2012 ACM Subject Classification Software and its engineering → Software notations and tools;
Computing methodologies → Natural language processing

Keywords and phrases issue classification, retrieval-augmented generation, fine-tuning, large language models, mining software repositories

Digital Object Identifier 10.4230/LIPICs...

1 Introduction

2 Related Work

Early IRC studies employed data mining techniques and ML classifiers [2, 20, 14]. For example, Antoniol et al. [2] leveraged a naive Bayes classifier to distinguish bugs from other issue reports. However, these methods showed low performance due to their limited semantic understanding of the issue report text [15].

To overcome this, recent methods have adopted transformer models due to their strong contextual understanding [7, 5, 26, 18, 19]. These studies typically fine-tune the models on historical issues in a supervised setting. For instance, Izadi et al. [18] fine-tuned a RoBERTa [23] model to classify issues into *Bugs*, *Features*, or *Questions*. Similarly, Trautsch et al. [26] trained a seBERT [27] model to classify issues into the same three categories. Although transformer-based approaches achieve strong results, they rely on large amounts of project-specific data to be effective, hindering their generalizability on projects with a small number of labeled issues [15, 18]. Additionally, they struggle to accurately classify minority classes with fewer training examples [19, 7].

These challenges have motivated the exploration of LLMs, which exhibit strong text classification capabilities thanks to extensive pretraining on massive corpora. [15, 4, 3, 6, 28, 10]. Studies show that LLMs perform well in zero-shot settings [10, 3]. SOTA LLM-based methods achieved the best performance by fine-tuning the models on labeled datasets [15, 4]. For example, Aracena et al. [4] instruction-tuned several models on the NLBSE '23 and '24 datasets.

3 Approach

3.1 Problem Formulation

Given a query issue report, we classify it into only one of the three labels: *bug*, *feature*, or *question*. Only the issue's title and description text are used for classification with no other signals from the issue tracker. The defined label space is adopted directly from prior IRC studies [3, 4, 18, 20, 2].



© Anonymous author(s);

licensed under Creative Commons License CC-BY 4.0

Leibniz International Proceedings in Informatics

LIPICs Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

42 **3.2 Indexing and Retrieval**

43 To prepare the issues for retrieval based on textual similarity, we first embed their textual
 44 content (i.e., *title + description*). We do not embed the issue labels. Instead, we retain
 45 them as metadata associated with each embedded issue. To generate the embeddings, we use
 46 the `all-MiniLM-L6-v2` model from the Sentence-Transformers library [24]. We selected this
 47 model due to its lightweight architecture, fast performance, and effectiveness on issue report
 48 text [17]. The embeddings are then indexed using the Faiss vector database library [12],
 49 which provides efficient high-dimensional similarity retrieval. At the retrieval stage, the query
 50 issue is first embedded with the same model. Next, Faiss calculates the cosine similarity
 51 of indexed issue reports to the query, and returns the top-K nearest neighbors, ordered by
 52 descending similarity.

53 **3.3 RAGTAG**

54 RAGTAG is our proposed retrieval-augmented few-shot prompting approach for IRC. RAG-
 55 TAG uses the retrieval pipeline described in Section 3.2 to retrieve the top-K nearest neighbors
 56 to the query issue, and present them as classified examples to the LLM in a few-shot prompt.
 57 This enables a more accurate classification, based on in-context learning [6]. The choice to
 58 present the top-K nearest neighbors as examples is based on two insights from prior research:
 59 First, retrieval-based example selection outperforms random selection [21]. Second, the most
 60 similar neighbors tend to share the same labels [11]; therefore, they are informative for the
 61 classification.

62 We create the few-shot prompt in three parts. First, the system message frames the
 63 classification task and the label space. Second, the top-K nearest neighbors are listed in
 64 retrieval order (descending similarity), each formatted as *title + description + label*. Finally,
 65 the query issue is appended to the end of the prompt, ensuring it receives high attention
 66 from the LLM [22]. $K = 0$ produces a zero-shot prompting baseline, similar to prior work [8].

67 We evaluate RAGTAG at top-K values in $\{0, 1, 3, 6, 9\}$. Details for the K value grid
 68 selection are provided in Section 5.1. Since the median prompt size at $K = 9$ in our dataset
 69 is $\sim 6,300$ tokens, the maximum sequence length is set to 8,192 tokens. Roughly 30% of the
 70 context quota is reserved for the query issue and 70% for the neighbors. If the neighbors
 71 exceed their overall quota, they are truncated proportionally based on their length (longer
 72 neighbors get more truncation). The query issue gets truncated if it exceeds its budget as
 73 well. Truncation happens on the right side of the description text, preserving all titles and
 74 labels (if any).

75 The LLM is instructed to generate the predicted label wrapped in XML tags (e.g.
 76 `<label>bug</label>`). The labels of the presented examples also follow the same format.
 77 This technique enables easy detection of the predicted label from the LLM’s generation, and
 78 reduces variance in LLM outputs caused by prompt format sensitivity [25]. We use a regular
 79 expression to extract the predicted label. Unparseable outputs are marked as *invalid*.

80 **3.4 VOTAG**

81 VOTAG is our proposed non-parametric retrieval IRC method. VOTAG uses the retrieval
 82 pipeline described in Section 3.2 to retrieve the top-K nearest neighbors of the query issue
 83 report, then predicts the final label with a similarity-weighted vote on their labels. This
 84 approach follows the standard K-nearest-neighbor classification setup [11], with neighbors

85 weighted by their similarity to the query [13].¹ Formally, for a query issue q with retrieved
 86 neighbors $\{(n_i, l_i, s_i)\}_{i=1}^K$, where n_i is the i -th nearest neighbor with label l_i and cosine
 87 similarity s_i to q , the label is predicted by:

$$88 \quad \hat{y} = \arg \max_{c \in \mathcal{L}} \sum_{i=1}^K s_i \cdot \mathbf{1}[l_i = c], \quad (1)$$

89 where $\mathcal{L}$ is the label set and $\mathbf{1}[\cdot]$ is the indicator function. VOTAG breaks ties by returning
 90 the label of the most similar neighbor in the tied classes.

91 VOTAG has three main roles: First, it serves as a fast and cost-efficient baseline that
 92 LLM-based approaches must clear to justify their computational cost. Second, it is used
 93 as an ablation to distinguish the contributions of the LLM and retrieval in the RAGTAG
 94 method. Third, VOTAG’s performance as a function of K informs the K value grid evaluated
 95 for RAGTAG.

96 3.5 Fine-Tuning

97 SOTA IRC approaches have achieved strong performance by supervised fine-tuning of
 98 LLMs [15, 4, 3], using Low-Rank Adaptation (LoRA) [16]. In these studies, the model is
 99 trained on labeled issue reports and predicts the labels for unseen test issues. Each training
 100 example is formatted as an instruction prompt which includes the issue’s title and description
 101 as input, and its label as output. At inference, the trained model receives the test issue in
 102 the same format, and predicts the empty label field. Regular expressions are used to extract
 103 the prediction from the output, and unparseable outputs are marked as *invalid*. We use this
 104 established methodology as our fine-tuning baseline.

105 For memory efficiency, each model is loaded in 4-bit quantization with Unsloth². The
 106 LoRA rank and alpha are set to 16, with a learning rate of 2×10^{-4} and the paged AdamW
 107 8-bit optimizer. Training runs for 3 epochs with a per-device batch size of 1 and gradient
 108 accumulation steps of 16. The max sequence length is set to 2,048 tokens, which covers
 109 94.9% of the issues in our dataset. Issues that exceed this budget get truncated from the
 110 right side of their description text to fit.

111 4 Experimental Setup

112 4.1 Dataset and Settings

113 We use the dataset introduced in Heo et al. [15] for our experiments. The dataset has a total
 114 of 6,600 issue reports gathered from eleven active open-source projects: `ansible/ansible`,
 115 `bitcoin/bitcoin`, `dart-lang/sdk`, `dotnet/roslyn`, `facebook/react`, `flutter/flutter`,
 116 `kubernetes/kubernetes`, `microsoft/TypeScript`, `microsoft/vscode`, `opencv/opencv`, and
 117 `tensorflow/tensorflow`. Each project has 600 issue reports with an equal distribution
 118 across the labels (*bug*, *feature*, *question*). To support their LLM fine-tuning method, Heo et
 119 al. divided the dataset into 50/50 train and test splits, stratified by project and label. This
 120 results in 300 train and 300 test issue reports per project.

121 We also adopt Heo et al.’s project-specific (PS) and project-agnostic (PA) configurations
 122 to examine how data scope affects IRC performance. In the PS setting, each project’s train

¹ Squared-similarity weighting and uniform majority were also evaluated, but were outperformed by similarity-weighted voting.

² <https://unsloth.ai>

and test data are used in isolation, resulting in eleven different train-test pairs. In the PA setting, the train splits of all projects are merged into a single training set of 3,300 issue reports. In both configurations, RAGTAG and VOTAG only use the training splits as their retrieval index to ensure fair comparison with the fine-tuning baseline.

4.2 Models

We use the Qwen2.5-Instruct model family [29] in our evaluations with four different parameter sizes: 3B, 7B, 14B, and 32B. Qwen2.5 is an open-source LLM with strong performance on general-purpose benchmarks. Using the same model at different parameter sizes isolates the effect of model scale from model architecture across LLM-based IRC approaches. Model temperature is set at 0.1 for deterministic classifications.

4.3 Evaluation Metrics

We report macro-averaged precision, recall, and F_1 over the three labels (**bug**, **feature**, **question**), along with overall accuracy. Per-class metrics are reported when the asymmetry across classes is itself the subject of analysis. Outputs that fail to parse a valid label (after the XML-tag extraction and regex fallback described in ??) are counted as incorrect, and an *invalid rate* is reported separately so any parsing-induced loss is auditable.

4.3.0.1 Aggregation convention (PA vs PS).


The PA setting produces a single set of predictions over the full 3,300-issue test set, and we compute macro F_1 directly over that set. The PS setting produces eleven sets of predictions, one per project, each over its own ~ 300 -issue test slice. Throughout the paper we report PS results using *pooled* aggregation: the eleven per-project prediction sets are concatenated into one 3,300-issue test set and macro F_1 is computed once on the pool. This treats each test issue as one observation in both settings and yields apples-to-apples comparisons between PA and PS. The alternative — computing macro F_1 separately within each project and averaging the eleven numbers — weights projects rather than issues; because macro F_1 is non-linear in the underlying confusion-matrix counts, the two conventions disagree, and on our data the per-project mean is consistently 0.002–0.011 macro F_1 below the pooled value with the gap largest for fine-tuning. We adopt pooling as the canonical reporting convention to remove this asymmetry; per-project breakdowns appear only where the project is the unit of analysis and are explicitly labelled as such.

4.4 Setup Hardware

5 Evaluations

5.1 RQ1) To what extent can similarity-weighted voting on the nearest neighbors (VOTAG) classify issue reports, and at what point does adding neighbors stop helping?

We evaluate VOTAG’s performance on both PA and PS settings at k values 1 to 30. The resulting macro F1 scores are shown in Figure 1. VOTAG’s best macro F1 is 0.595 ($k=15$) in PS and 0.604 ($k=16$) in PA. Both settings reach 99.8% (PS) and 98.5% (PA) of their respective peak by $k=7$. When adding neighbors beyond the peak, PS loses 0.015 and PA loses 0.011 by $k=30$.



■ **Figure 1 (left)** VOTAG macro F_1 as a function of k in both PS and PA settings. Circles mark peak performance. **(right)** Per-class F_1 at each setting’s best k ($k=15$ for PS, $k=16$ for PA).

PA consistently outperforms PS at every k . However, the performance gap is marginal: averaging 0.007 macro F_1 , with a maximum of 0.015 at $k=18$. This similar performance is due to the large overlap of the top- k neighbors in both settings. Specifically, PA and PS share 84–90% of the top- k neighbors across $k \in [3, 30]$. Therefore, even when using the full 3,300-issue retrieval index, most of the retrieved neighbors are from the same project. This means issue reports from the same project tend to be closer in the embedding space.

Question consistently has the weakest performance among the three labels at every k in both settings. At best k , PS per-class F_1 is 0.630 for bug, 0.614 for feature, and 0.541 for question. The corresponding PA values are 0.632, 0.629, and 0.550 (Figure 1). Additionally, questions are more frequently misclassified as bugs: 32% of question issues are predicted as bugs in both settings, while 18% (PS) and 17% (PA) are predicted as features. Since VOTAG predicts the label with the most accumulated similarity weight, this misclassification asymmetry indicates that question issues are more similar to bugs than to features in the embedding space.


5.2 RAGTAG

We evaluate RAGTAG at k values $\{0, 1, 3, 6, 9, 12, 15\}$ on both PA and PS settings, where $k=0$ is the zero-shot baseline [8]. k value grid selection is informed by VOTAG’s peak performance around $k=15$ in the previous section. Figure 2 shows the macro F_1 of every model across this grid for both settings.

PS marginally outperforms PA at almost every (model, $k \geq 1$) pair evaluated (zero-shot is identical between settings since no retrieval is performed). The close performance gap is due to the same top- k neighbor overlap observed in the VOTAG analysis. The marginal improvement is significant because PS uses $11 \times$ less retrieval data per project: each PS index covers only its own 300 retrieval issues, while PA uses the full 3,300 across all 11 projects. Given this practical advantage, we conduct the rest of our RAGTAG analysis on the PS setting and report PS numbers throughout this section.

The best k grows with model size: Qwen-3B peaks at $k=3$ (macro $F_1 = 0.697$), Qwen-7B at $k=6$ (0.718), and both Qwen-14B and Qwen-32B at $k=12$ (0.732 and 0.767 respectively). This shows that larger models can use the context from additional examples more effectively. However, adding neighbors past $k=12$ has no benefit and slightly degrades performance: every model loses 0.001–0.015 macro F_1 going from $k=12$ to $k=15$.

Every RAGTAG configuration with $k \geq 1$ outperforms the zero-shot baseline. At each




■ **Figure 2** RAGTAG macro F_1 as a function of k for all four Qwen sizes. $k=0$ is the zero-shot baseline. The dotted horizontal line marks VOTAG’s best (0.604, PA $k=16$).

195 model’s peak, the gap over zero-shot averages $+0.076$ macro F_1 and ranges from $+0.055$
 196 (Qwen-7B) to $+0.087$ (Qwen-14B). This shows that LLMs benefit substantially from the
 197 most similar labeled examples for IRC.

198 Every RAGTAG configuration outperforms VOTAG’s best result (0.604 at PA $k=16$),
 199 with the smallest margin from Qwen-3B zero-shot ($F_1 = 0.613$, $+0.009$) and the largest from
 200 Qwen-32B PS at $k=12$ ($F_1 = 0.767$, $+0.163$). These results show that adding the LLM
 201 reasoning improves IRC on top of pure retrieval.

202 Similar to VOTAG, question consistently has the weakest performance among the three
 203 labels across all RAGTAG settings. Figure 3 shows per-class F_1 at each model’s best k .
 204 Even the smallest gaps between question and the other labels, observed at Qwen-32B with
 205 $k=12$, are noticeable: question is 0.060 macro F_1 below bug and 0.143 below feature. Similar
 206 to the VOTAG observations, questions tend to be more frequently misclassified as bugs.
 207 At zero-shot, across all four model sizes, 46–59% of true-question issues are predicted as
 208 bugs while only 5–16% are predicted as features. RAGTAG narrows this misclassification
 209 asymmetry but does not eliminate it: at each model’s best k , 26–36% of question issues are
 210 still predicted as bugs and only 9–15% as features.

211 These results show that the LLMs are already prone to misclassifying question issues as
 212 bugs, even before the top- k examples are used. Therefore, presenting bug-labeled examples
 213 for true-question query issues likely helps the misclassification. Based on this intuition, we
 214 propose a more balanced retrieval technique as an intervention to reduce question→bug
 215 misclassifications. To prevent damaging performance on true-bug issues, we apply this
 216 neighbor balancing only when the query is suspected to be a question. We use the label
 217 distribution of the retrieved top- k neighbors as the suspicion signal. Averaged across
 218 $k \in [1, 30]$, question query issues retrieve roughly balanced bug and question examples (mean
 219 $N_{\text{bug}} - N_{\text{question}} = -1.4$). Therefore, we treat a query issue as a suspected question when the
 220 question count in the top- k is within a margin m of the bug count ($N_{\text{bug}} - N_{\text{question}} \leq m$).



■ **Figure 3** Per-class F_1 for each model size at its best RAGTAG-PS configuration.

221 We select $m=3$ empirically from $m \in \{1, 2, 3, 4, 5\}$. Averaged across $k \in [1, 30]$, $m=3$ fires for
 222 79% of true-question issues while firing for only 41% of true-bug issues. This trade-off is
 223 more favorable than other margin values. We further confirmed this choice with a margin
 224 sweep on Qwen-3B (the smallest model), where $m=3$ had the best performance. We refer to
 225 this method as **Balanced RAGTAG** (BRAGTAG).

226 5.3 Balanced RAGTAG


227 Since the PS setting produced the best results for RAGTAG, we use PS as the default for
 228 evaluating BRAGTAG on the same $k \in \{1, 3, 6, 9, 12, 15\}$ grid. Figure 4 shows macro F1 vs
 229 k for both methods across all four model sizes. BRAGTAG outperforms RAGTAG at every
 230 $k \geq 6$. At smaller $k \in \{1, 3\}$, BRAGTAG slightly underperforms by 0.001–0.010 macro F1,
 231 because the trigger fires too often and removes every retrieved example, practically reverting
 232 the query to zero-shot (40% of queries at $k=1$ and 18% at $k=3$).

233 BRAGTAG’s best k almost always increases compared to RAGTAG: Qwen-3B from
 234 $k=3$ to $k=6$, Qwen-7B from $k=6$ to $k=12$, Qwen-14B from $k=12$ to $k=15$, and Qwen-32B
 235 unchanged at $k=12$. When the trigger fires, BRAGTAG removes all bug examples from the
 236 top- k . Therefore, A higher initial k is needed to leave enough remaining context for the LLM
 237 to use. At each method’s best k , BRAGTAG outperforms RAGTAG on every model. The
 238 paired bootstrap 95% CI on the (BRAGTAG – RAGTAG) macro F_1 difference excludes
 239 zero in all four cells:⁴ Qwen-3B +0.017 [+0.003, +0.031], Qwen-7B +0.020 [+0.007, +0.032],
 240 Qwen-14B +0.024 [+0.014, +0.034], and Qwen-32B +0.015 [+0.007, +0.022]. Table 1 reports
 241 the full best configuration comparison.

242 The performance gain comes from an increase on the question and bug classes with
 243 minimal impact on feature. Figure 5 shows per-class F_1 at each method’s best k . Question
 244 F_1 improves by 0.030–0.068 across all four model sizes. Bug F_1 stays within 0.018 of
 245 RAGTAG, with a slight loss for Qwen-3B and slight gains for the larger models. Feature

³ The margin cannot be large given the small typical gap between bug and question neighbor counts for a true-question.

⁴ 1,000 paired bootstrap resamples per model.



■ **Figure 4** Macro F_1 as a function of k across all model sizes for both RAGTAG and BRAGTAG (project-specific setting)

■ **Table 1** RAGTAG vs BRAGTAG at each model’s best k (PS, pooled).

Model	Method	k^*	Macro F_1	F_1^{bug}	F_1^{feat}	F_1^{q}	Invalid rate
Qwen-3B	RAGTAG	3	0.697	0.711	0.803	0.577	0.2%
	BRAGTAG	6	0.714	0.693	0.803	0.645	1.8%
Qwen-7B	RAGTAG	6	0.718	0.731	0.816	0.605	3.5%
	BRAGTAG	12	0.738	0.742	0.805	0.666	3.8%
Qwen-14B	RAGTAG	12	0.732	0.735	0.838	0.622	4.5%
	BRAGTAG	15	0.756	0.745	0.840	0.683	4.7%
Qwen-32B	RAGTAG	12	0.767	0.759	0.842	0.699	4.6%
	BRAGTAG	12	0.781	0.775	0.840	0.729	4.0%

F_1 is essentially unchanged. The question→bug misclassification rate also drops sharply: at each model’s best k , RAGTAG misclassifies 26–36% of true-question issues as bugs while BRAGTAG reduces this to 19–27%. Averaged across $k \in [1, 15]$, the rate drops from 31–40% (RAGTAG) to 24–36% (BRAGTAG). Question→feature misclassification is unaffected. This confirms that BRAGTAG’s improvement comes from the targeted reduction of question→bug misclassification.

Outputs from both approaches that failed to parse to a valid label are marked as *invalid* and counted as incorrect. Averaged across $k \in [1, 15]$, BRAGTAG has a lower invalid rate than RAGTAG (2.3% vs 3.0% across the four models).

5.4 Fine-Tuning

We fine-tune all models in both PA and PS settings. Figure 6(a) compares macro F_1 for fine-tuning in both settings against RAGTAG’s PS at each model’s best k . Fine-tune PA outperforms PS at every model size, by +0.033 (Qwen-3B), +0.084 (Qwen-7B), +0.082 (Qwen-14B), and +0.031 (Qwen-32B) macro F_1 . This is the opposite of what we observed for RAGTAG, where PS marginally outperformed PA. The 300 labeled issues per project




Figure 5 Per-class F_1 at each method’s best k across all model sizes for RAGTAG and BRAGTAG.




Figure 6 (a) Macro F_1 for RAGTAG at its best k in PS vs Fine-Tune in PS and PA across four model sizes. **(b)** Per-class F_1 gain from Fine-Tune PS to Fine-Tune PA ($\Delta F_1 = \text{PA} - \text{PS}$) for each class and model size.

exposed to the model in fine-tune PS is not sufficient to learn the label distributions. This trend matches prior findings on LLM fine-tuning for IRC [15].

RAGTAG’s PS outperforms Fine-Tune’s PS at every model size, by +0.021 (Qwen-3B) to +0.040 (Qwen-7B) macro F_1 . This shows that with only 300 labeled issues per project, retrieval-augmented few-shot is more effective than LoRA fine-tuning for IRC.

Figure 6(b) shows where the additional PA training data helps most in fine-tuning. Averaged across all four models, bug F_1 gains +0.070 going from PS to PA, question gains +0.053, and feature gains +0.050. Qwen-32B is the only exception to the PS→PA gain: its question F_1 drops by 0.010 while bug gains +0.070.

Fine-tune produces fewer invalid outputs than RAGTAG and BRAGTAG in both settings, because the model learns the output format during training. Averaged across the four models, fine-tune’s invalid rate is 0.28% in PA and 0.38% in PS, below RAGTAG’s 3.0% and BRAGTAG’s 2.3% (averaged across $k \in [1, 15]$).

5.5 Method Comparison

In a standard software development workflow, every issue needs a label, so any IRC approach must produce a valid output. However, a known limitation of LLMs is generating outputs that do not comply with the instructions [25, 1]. Our previous analysis also shows that all

■ **Table 2** RAGTAG, BRAGTAG, and Fine-Tune at each method’s best configuration with the scope-matched VOTAG fallback applied (VOTAG-PS for RAGTAG/BRAGTAG, VOTAG-PA for Fine-Tune). Train and Inference are wall-times on a single GPU (RAGTAG/BRAGTAG have no training phase). Total is their sum and excludes model load. RAM is the peak GPU memory observed. Pooled.

Model	Method	k^*	Macro F_1	F_1^{bug}	F_1^{feat}	F_1^{q}	RAM (GB)	Train (h)	Infer (h)	Total (h)
Qwen-3B	RAGTAG	3	0.696	0.710	0.801	0.577	2.9	–	0.18	0.18
	BRAGTAG	6	0.720	0.707	0.803	0.652	2.9	–	0.22	0.22
	Fine-Tune	–	0.711	0.739	0.789	0.604	5.5	0.31	0.11	0.42
Qwen-7B	RAGTAG	6	0.729	0.751	0.815	0.620	6.8	–	0.50	0.50
	BRAGTAG	12	0.751	0.764	0.806	0.682	6.8	–	0.60	0.60
	Fine-Tune	–	0.762	0.761	0.848	0.677	9.9	0.48	0.10	0.58
Qwen-14B	RAGTAG	12	0.746	0.757	0.838	0.642	12.6	–	1.37	1.37
	BRAGTAG	15	0.771	0.773	0.841	0.699	12.6	–	1.35	1.35
	Fine-Tune	–	0.785	0.792	0.828	0.736	16.7	1.49	0.22	1.71
Qwen-32B	RAGTAG	12	0.779	0.780	0.842	0.715	22.3	–	4.62	4.62
	BRAGTAG	12	0.792	0.795	0.840	0.743	22.3	–	4.15	4.15
	Fine-Tune	–	0.771	0.791	0.853	0.669	31.5	2.28	0.44	2.71

three LLM-based approaches produce some invalid outputs, even fine-tune at 0.28–0.38%. To always have an output, we propose falling back to VOTAG for invalid LLM predictions. VOTAG never produces an invalid label and adds no LLM inference cost.

We use this fallback to compare RAGTAG, BRAGTAG, and fine-tune at their best configurations established by our experiments: RAGTAG and BRAGTAG at their best k on PS, and fine-tune on PA. Using the same fallback technique across all three ensures a fair comparison at each method’s peak performance. To match each method’s best data scope, we use VOTAG PS as the fallback for RAGTAG and BRAGTAG, and VOTAG PA for fine-tune.

Table 2 reports macro and per-class F_1 at each method’s best configuration with the fallback applied. RAGTAG has competitive macro F_1 with fine-tuning: aggregated across the four model sizes (13,200 paired predictions), fine-tune leads RAGTAG by only 0.020 macro F_1 (paired bootstrap 95% CI $[-0.028, -0.013]$). RAGTAG is closer to fine-tune at the smallest and largest models and farther in the middle, with per-model (RAGTAG – fine-tune) macro F_1 deltas (paired bootstrap 95% CIs in brackets) of -0.015 $[-0.032, +0.002]$ (Qwen-3B), -0.033 $[-0.047, -0.017]$ (Qwen-7B), -0.040 $[-0.054, -0.025]$ (Qwen-14B), and $+0.008$ $[-0.008, +0.022]$ (Qwen-32B). RAGTAG and fine-tune are statistically tied at Qwen-32B.

BRAGTAG’s balanced retrieval closes this remaining gap to statistical equivalence with fine-tune: the aggregate (BRAGTAG – fine-tune) difference is $+0.001$ with paired bootstrap 95% CI $[-0.007, +0.008]$, passing TOST equivalence at $\delta=0.01$.⁵ Similar to RAGTAG, BRAGTAG performs better relative to fine-tune at the smallest and largest models: BRAGTAG has a slight edge over fine-tune at Qwen-3B and Qwen-32B while fine-tune leads on the mid-sized ones, with per-model (BRAGTAG – fine-tune) deltas of $+0.009$ $[-0.008, +0.028]$ (Qwen-3B), -0.011 $[-0.026, +0.004]$ (Qwen-7B), -0.014 $[-0.030, +0.001]$ (Qwen-14B), and $+0.021$ $[+0.006, +0.035]$ (Qwen-32B).

BRAGTAG has the best question F_1 on all models except Qwen-14B. The lead for bug and feature classes alternates between Fine-tune and BRAGTAG each winning two of four models per class. These results show that BRAGTAG is able to compete with fine-tune

⁵ 1,000 paired bootstrap resamples per model for both RAGTAG vs fine-tune and BRAGTAG vs fine-tune comparisons.




Figure 7 (a) Peak GPU RAM by model size: RAGTAG/BRAGTAG inference vs Fine-Tune training + inference. (b) Total GPU time by model size: RAGTAG and BRAGTAG vs Fine-Tune

due to its targeted question→bug correction, boosting question F_1 without sacrificing bug or feature. Averaged across the four models, BRAGTAG is also the most class-balanced approach, with the lowest per-class F_1 standard deviation (0.053 vs 0.066 for fine-tune and 0.076 for RAGTAG) and the highest worst-class F_1 floor (0.694 vs 0.672 and 0.639).

The methods have different hardware requirements. Figure 7(a) shows peak GPU RAM by model. RAGTAG and BRAGTAG require identical GPU memory that grows with model size, while fine-tune requires more at every size due to training overheads such as, gradients, optimizer state, and activations retained for the backward pass. Averaged across the model sizes, RAGTAG and BRAGTAG take 33% less GPU memory than fine-tune (47%, 31%, 25%, and 29% less at Qwen-3B, 7B, 14B, and 32B respectively).

Figure 7(b) shows total GPU time at each model size. Fine-tune takes longer than RAGTAG at Qwen-3B, 7B, and 14B. But at Qwen-32B fine-tune (2.71 h) takes less time than both RAGTAG (4.62 h) and BRAGTAG (4.15 h). Inference time on each model naturally grows with prompt length. Fine-tune processes only the query issue, while RAGTAG and BRAGTAG process significantly more context due to the additional few-shot examples per query. As a result, fine-tune inference takes much less time than rag-based inference, and fine-tune’s total run-time is dominated by training. At Qwen-7B and Qwen-32B, BRAGTAG’s inference cost alone exceeds fine-tune’s combined training and inference time. BRAGTAG’s intervention shortens the average context enough at Qwen-14B and Qwen-32B for BRAGTAG to run less than RAGTAG. However, BRAGTAG runs longer at Qwen-3B and Qwen-7B because the intervention could not compensate for the higher best k value.

Notably, RAGTAG and BRAGTAG in PS retrieve only from 300 labeled issues per project, while fine-tune PA trains on the full 3,300 issues across all 11 projects to reach its best configuration. Therefore, at $11\times$ less labeled data per project, RAGTAG manages to be competitive with fine-tune, and BRAGTAG fully closes the remaining gap to statistical equivalence.

6 Discussion

6.1 Failure Analysis

In this section we conduct a manual qualitative analysis to better understand the misclassification patterns across the explored IRC methods. We inspect the misclassifications produced by Qwen-32B (the strongest model) under each method’s best configuration. From all the

misclassifications across RAGTAG, BRAGTAG, and fine-tune, we drew a random sample of 50 misclassifications per method (150 total). Each Method’s sample includes 30 true-question, 10 true-bug, and 10 true-feature errors. Each case was analyzed by two authors independently, with discrepancies resolved through discussions. The identified failure categories and their frequencies are as follows:

Wrong template format (77.33%). Issue reports whose content and ground-truth label contradicts their chosen template and structure. All of these errors are from questions and feature requests that the user filed under a bug-report template (*Steps to reproduce, Expected, Actual, System Information*). Consequently, the structure of the bug template misleads the LLM classifier towards predicting bug, regardless of the issue’s content. For example, `microsoft/vscode#193985`⁶ is an actual support question that opens with a literal **Type:** Bug because it uses VS Code’s bug template.

Hybrid intent (13.33%). Issues that could have multiple defensible labels due to their combined intents. These errors mostly take place when a fix or a feature request is layered in a question or bug. For example, in `kubernetes/kubernetes#120364`⁷ the author both reports a concrete bug and in the same body explicitly requests a new feature: “*Please, give us multiarchitecture image! and fix monoarchitecture images too!*”

Label noise (9.33%). Issues whose ground-truth label does not match its content. For example, `tensorflow/tensorflow#61710`⁸ is titled “*Will there a MLP model in the future version?*” which requests a new API feature for building MLPs. Yet the dataset labels it as a bug.

We also sampled 30 issues with *invalid* predictions per method (60 total) and observed two failure modes. In roughly 74% of cases, the model kept generating prose past the label slot — either continuing the issue-body content (extra *Steps to reproduce* fields, stack traces, code) or its own chain-of-thought reasoning — without ever emitting a structured label. The remaining 26% are degenerate hallucinations: repeated digits, repeated punctuation, or out-of-schema labels (e.g., `docs`) that cannot be parsed into any of the three classes.

6.2 Key Insights and Practical Implications

The experiments and the failure analysis in Section 6.1 together yield several actionable insights for practitioners deploying IRC and for researchers building on this benchmark.

The bug bias is structural, not parametric. The misclassification asymmetry of question issues toward bug appears in every method we evaluated. VOTAG, which has no LLM, predicts bug for 32% of true-question issues at its best k (Section 5.1). Zero-shot RAGTAG, which has no retrieval, predicts bug for 46–59% of true-question issues across the four model sizes (Section 5.2). Fine-tuning on the same train split inherits the skew as well, with question consistently the lowest- F_1 class across all three approaches (Table 2). The bias survives both the removal of the LLM and the removal of retrieval, and prior work reports the same dominant error pattern on different LLMs and datasets [8, 9] — which suggests that the asymmetry is a property of the IRC task and the embedding space rather than of any specific model family. A direct measurement on our data supports this reading: for true-question queries, the average label imbalance among the top- k neighbors is $N_{\text{bug}} - N_{\text{question}} = -1.4$ averaged across $k \in [1, 30]$, while bug queries retrieve overwhelmingly

⁶ <https://github.com/microsoft/vscode/issues/193985>

⁷ <https://github.com/kubernetes/kubernetes/issues/120364>

⁸ <https://github.com/tensorflow/tensorflow/issues/61710>

bug neighbors. BRAGTAG targets the retrieval side of this mechanism by removing bug examples from the prompt only when the query’s neighborhood is roughly balanced between bug and question, and the failure analysis confirms the intervention reaches its target population: the retrieval-skew rate drops from 26% (RAGTAG) to 14% (BRAGTAG), exactly halving the share of misclassifications driven by a wrong-class-dominated neighborhood.

Statistical equivalence at the dataset-noise ceiling. The aggregate (BRAGTAG – fine-tune) macro F_1 difference is +0.001 with paired bootstrap 95% CI $[-0.007, +0.008]$, passing TOST equivalence at $\delta=0.01$ (Section 5.5), despite BRAGTAG retrieving from only 300 labeled issues per project while fine-tune trains on 3,300. The failure analysis recasts this equivalence: only 3% of residual misclassifications are clean errors with no template confound and no plausible label ambiguity (Section 6.1). The other 97% involve either a wrong-template confound (76%, where the bug-report scaffolding pushes the classifier away from the user’s true intent), label noise (9%), or genuinely ambiguous hybrid intents (13%). Both methods do not match because they are independently equally accurate; both reach a regime where the residual error is dominated by inputs whose templates or labels make the classification ambiguous within the three-class schema. The practical implication is that further gains on this benchmark are unlikely to come from stronger LLMs alone — they will require either label cleaning, template-aware preprocessing, or a richer label space.

Match the method to the labeled-data scope. RAGTAG and fine-tuning respond to the PS–PA distinction in opposite directions. RAGTAG’s best setting is PS at every (model, $k \geq 1$) pair (Section 5.2); fine-tune’s best setting is PA at every model size, by +0.031 to +0.084 macro F_1 (Section 5.4). The structural reason is that PS retrieval indexes 300 issues from the same project as the query, and same-project issues cluster together in the embedding space, so the small index covers most of the relevant neighborhood; we measured PA and PS to share 84–90% of the top- k neighbors across $k \in [3, 30]$ in Section 5.1. PS is therefore not just a smaller retrieval pool but a near-complete one for a retrieval-based method. Fine-tuning has no such property: 300 labeled issues per project are insufficient to fit the question/bug/feature distinction, and the model benefits from the broader 3,300-issue PA pool. The practical recommendation is therefore conditional on what labeled data a project has available: with ~ 300 in-project labeled issues, retrieval-augmented few-shot is the stronger choice, while $\sim 3,300$ issues spanning multiple projects make fine-tuning competitive. BRAGTAG extends the first regime: at 300 in-project labels, it matches fine-tune’s 3,300-issue PA performance to within statistical equivalence. The $11\times$ data factor between settings is the operationally important takeaway for projects without a cross-project label corpus.

Cost crossover at large model size. Wall-time and memory characteristics differ qualitatively across methods. Fine-tune carries a fixed training cost that grows roughly linearly with model size, while RAGTAG and BRAGTAG have no training phase but scale inference time with prompt length. The ordering inverts at the largest model size: at Qwen-32B, fine-tune’s total wall-time (2.71 h) is shorter than both RAGTAG (4.62 h) and BRAGTAG (4.15 h), even though fine-tune is more expensive at every smaller size. The crossover is driven by long-context inference cost overtaking training cost as model size grows but the best k stays fixed at 12 (Section 5.5). The retrieval-based methods are therefore most cost-favorable in the model-size range a single consumer GPU can host (3B, 7B, 14B); at 32B and beyond, fine-tuning may be the more efficient workflow if the labeled-data conditions favor it.

Class balance is a hidden axis of comparison. BRAGTAG is the most class-balanced of the three methods. Averaged across the four model sizes, the per-class F_1 standard deviation is 0.053 for BRAGTAG, 0.066 for fine-tune, and 0.076 for RAGTAG, and

the worst-class F_1 floor is 0.694, 0.672, and 0.638 respectively (Table 2). Aggregate macro F_1 does not capture this: a method that is uniformly competent across the three labels can be preferable for triage deployment to one with the same macro F_1 but a noticeable per-class trough. BRAGTAG’s intervention specifically lifts the question floor without sacrificing the other two classes, making it the most stable choice when no single label can be allowed to degrade.

Fine-tuning memorizes template surface signals. A finding that emerges only from the qualitative analysis is that fine-tune’s residual errors include a notable concentration on a single tooling-driven template artifact: issues filed via VS Code’s “*Help* → *Report Issue*” menu, which embeds a literal “*Type: Bug*” marker regardless of intent (Section 6.1). Fine-tuning learned this surface signal as bug-predictive and over-applies it. Retrieval-based methods, which never see all 3,300 issues at once, are less prone to memorizing such template artifacts. This is a concrete, dataset-grounded reason to prefer in-context methods when a project’s issue tracker contains tooling-driven template noise.

6.3 Future Work

Tighter debiasing interventions. BRAGTAG’s margin-based filter halves the retrieval-skew rate (26%→14%) but does not eliminate it. The remaining 14% of misclassifications still occur on queries whose top- k neighborhood is dominated by the wrong class even after bug examples are removed, suggesting room for a tighter intervention. Two directions follow naturally from the analysis: extending the trigger to feature-bias as well (the current rule only addresses the question-vs-bug imbalance), and replacing the discrete margin filter with a continuous embedding-space rebalancing that re-weights neighbors by class frequency rather than dropping them. Either direction is implementable as a drop-in replacement for our `_debias_neighbors` function and would preserve BRAGTAG’s deployment-friendly properties (no training, no extra inference cost beyond the retrieval the method already performs).

Label cleaning and richer label spaces. The combined share of label noise (9%), hybrid intents (13%), and wrong template format (76%) leaves only 3% of residual errors as clean model failures (Section 6.1), bounding the headroom on this benchmark for any classifier within the three-class $\{bug, feature, question\}$ schema. Two complementary follow-ups suggest themselves. The first is template-aware preprocessing that strips or normalizes bug-template scaffolding (e.g., removing “*Type: Bug*” auto-markers and *Steps to reproduce* headers) so the classifier sees the user’s actual content rather than the tooling’s framing. The second is a richer label space that can express the patterns the failure analysis surfaced as recurring: a “*support-with-repro*” label for bug-template-shaped support requests and an explicit “*hybrid*” label for issues that legitimately combine two intents. Both directions would also benefit fine-tuning by removing systematic noise that a parametric model otherwise risks memorizing as signal.

Generalization to other labeled-data conditions and platforms. Our analysis covers an 11-project GitHub benchmark with a balanced 600 issues per project, which is a relatively favorable condition for retrieval. Two extensions would test the boundary of the regimes identified in Section 6.2. First, evaluating on imbalanced label distributions (the more typical case in industrial issue trackers) would test whether the retrieval-misled-error mechanism scales differently when one label dominates the index. Second, evaluating on issue trackers other than GitHub (e.g., Jira, Bugzilla) would test whether the wrong-template-format pattern generalizes; we expect platforms with stricter label vocabularies to show less of this confound but more of a different kind, such as field-misuse or sub-label drift. Both

474 directions could reuse the methodology and replication package released here.

475 **7 Threats to Validity**

476 **8 Conclusion**

477 **9 Data Availability**

478 **References**

- 479 **1** Arshia Akhavan, Alireza Hosseinpour, Abbas Heydarnoori, and Mehdi Keshani. LinkAn-
480 chor: An autonomous LLM-based agent for issue-to-commit link recovery. *arXiv preprint*
481 *arXiv:2508.12232*, 2025.
- 482 **2** Giuliano Antoniol, Kamel Ayari, Massimiliano Di Penta, Foutse Khomh, and Yann-Gaël
483 Guéhéneuc. Is it a bug or an enhancement? a text-based approach to classify change requests.
484 In *Proceedings of the 2008 conference of the center for advanced studies on collaborative*
485 *research: meeting of minds*, pages 304–318, 2008.
- 486 **3** Gabriel Aracena, Kyle Luster, Fabio Santos, Igor Steinmacher, and Marco A. Gerosa. Applying
487 large language models api to issue classification problem, 2024. URL: [https://arxiv.org/](https://arxiv.org/abs/2401.04637)
488 [abs/2401.04637](https://arxiv.org/abs/2401.04637), *arXiv:2401.04637*.
- 489 **4** Gabriel Aracena, Kyle Luster, Fabio Santos, Igor Steinmacher, and Marco A Gerosa. Applying
490 large language models to issue classification: Revisiting with extended data and new models.
491 *Science of Computer Programming*, page 103333, 2025.
- 492 **5** Shikhar Bharadwaj and Tushar Kadam. Github issue classification using bert-style models. In
493 *Proceedings of the 1st International Workshop on Natural Language-based Software Engineering*,
494 pages 40–43, 2022.
- 495 **6** Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla
496 Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language
497 models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901,
498 2020.
- 499 **7** Giuseppe Colavito, Filippo Lanubile, and Nicole Novielli. Issue report classification using
500 pre-trained language models. In *Proceedings of the 1st International Workshop on Natural*
501 *Language-based Software Engineering*, pages 29–32, 2022.
- 502 **8** Giuseppe Colavito, Filippo Lanubile, and Nicole Novielli. Benchmarking large language models
503 for automated labeling: The case of issue report classification. *Information and Software*
504 *Technology*, page 107758, 2025.
- 505 **9** Giuseppe Colavito, Filippo Lanubile, Nicole Novielli, and Luigi Quaranta. Impact of data
506 quality for automatic issue classification using pre-trained language models. *Journal of Systems*
507 *and Software*, 210:111838, 2024.
- 508 **10** Giuseppe Colavito, Filippo Lanubile, Nicole Novielli, and Luigi Quaranta. Leveraging gpt-like
509 llms to automate issue labeling. In *Proceedings of the 21st International Conference on Mining*
510 *Software Repositories*, pages 469–480, 2024.
- 511 **11** Thomas Cover and Peter Hart. Nearest neighbor pattern classification. *IEEE transactions on*
512 *information theory*, 13(1):21–27, 1967.
- 513 **12** Matthijs Douze, Alexandr Guzhva, Chengqi Deng, Jeff Johnson, Gergely Szilvasy, Pierre-
514 Emmanuel Mazaré, Maria Lomeli, Lucas Hosseini, and Hervé Jégou. The faiss library. *IEEE*
515 *Transactions on Big Data*, 2025.
- 516 **13** Sahibsingh A Dudani. The distance-weighted k-nearest-neighbor rule. *IEEE Transactions on*
517 *Systems, Man, and Cybernetics*, (4):325–327, 1976.
- 518 **14** Qiang Fan, Yue Yu, Gang Yin, Tao Wang, and Huaimin Wang. Where is the road for issue
519 reports classification based on text mining? In *2017 ACM/IEEE International Symposium on*
520 *Empirical Software Engineering and Measurement (ESEM)*, pages 121–130. IEEE, 2017.

- 521 **15** Jueun Heo and Seonah Lee. A study on applying large language models to issue classification.
522 In *2025 IEEE/ACM 33rd International Conference on Program Comprehension (ICPC)*, pages
523 1–11. IEEE Computer Society, 2025.
- 524 **16** Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuezhi Li, Shean Wang, Liang
525 Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *Iclr*, 1(2):3,
526 2022.
- 527 **17** Huihui Huang, Ratnadira Widyasari, Ting Zhang, Ivana Clairine Irsan, Jieke Shi, Han Wei
528 Ang, Frank Liauw, Eng Lieh Ouh, Lwin Khin Shar, Hong Jin Kang, et al. Back to the basics:
529 Rethinking issue-commit linking with llm-assisted retrieval. *arXiv preprint arXiv:2507.09199*,
530 2025.
- 531 **18** Maliheh Izadi. Catiss: An intelligent tool for categorizing issues reports using transformers. In
532 *Proceedings of the 1st international workshop on natural language-based software engineering*,
533 pages 44–47, 2022.
- 534 **19** Maliheh Izadi, Kiana Akbari, and Abbas Heydarnoori. Predicting the objective and priority
535 of issue reports in software repositories. *Empirical Software Engineering*, 27(2):50, 2022.
- 536 **20** Rafael Kallis, Andrea Di Sorbo, Gerardo Canfora, and Sebastiano Panichella. Ticket tagger:
537 Machine learning driven issue classification. In *2019 IEEE International Conference on*
538 *Software Maintenance and Evolution (ICSME)*, pages 406–409. IEEE, 2019.
- 539 **21** Jiachang Liu, Dinghan Shen, Yizhe Zhang, William B Dolan, Lawrence Carin, and Weizhu
540 Chen. What makes good in-context examples for gpt-3? In *Proceedings of Deep Learning*
541 *Inside Out (DeeLIO 2022): The 3rd workshop on knowledge extraction and integration for*
542 *deep learning architectures*, pages 100–114, 2022.
- 543 **22** Nelson F Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni,
544 and Percy Liang. Lost in the middle: How language models use long contexts. *Transactions*
545 *of the association for computational linguistics*, 12:157–173, 2024.
- 546 **23** Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy,
547 Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized bert
548 pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019.
- 549 **24** Nils Reimers and Iryna Gurevych. Sentence-bert: Sentence embeddings using siamese bert-
550 networks. *arXiv preprint arXiv:1908.10084*, 2019.
- 551 **25** Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. Quantifying language models’
552 sensitivity to spurious features in prompt design or: How i learned to start worrying about
553 prompt formatting. *arXiv preprint arXiv:2310.11324*, 2023.
- 554 **26** Alexander Trautsch and Steffen Herbold. Predicting issue types with sebert. In *Proceedings*
555 *of the 1st international workshop on natural language-based software engineering*, pages 37–39,
556 2022.
- 557 **27** Julian Von der Mosel, Alexander Trautsch, and Steffen Herbold. On the validity of pre-trained
558 transformers for natural language processing in the software engineering domain. *IEEE*
559 *Transactions on Software Engineering*, 49(4):1487–1507, 2022.
- 560 **28** Jason Wei, Maarten Bosma, Vincent Y Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan
561 Du, Andrew M Dai, and Quoc V Le. Finetuned language models are zero-shot learners. *arXiv*
562 *preprint arXiv:2109.01652*, 2021.
- 563 **29** An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan
564 Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2.5 technical report. *arXiv preprint*
565 *arXiv:2412.15115*, 2024.