

# TODO: Paper Title

Anonymous author

Anonymous affiliation

## Abstract

**Background.** TODO: 1–2 sentences on the problem and why it matters.

**Aims.** TODO: the research questions or specific objectives.

**Method.** TODO: dataset, models, approaches compared.

**Results.** TODO: headline numbers, ordered by RQ.

**Conclusions.** TODO: what this means and the takeaway for practitioners.

**2012 ACM Subject Classification** Software and its engineering → Software notations and tools;  
Computing methodologies → Natural language processing

**Keywords and phrases** issue classification, retrieval-augmented generation, fine-tuning, large language models, mining software repositories

**Digital Object Identifier** 10.4230/LIPICs...

## 1 Introduction

## 2 Related Work

Early IRC studies employed data mining techniques and ML classifiers [1, 18, 12]. For example, Antoniol et al. [1] leveraged a naive Bayes classifier to distinguish bugs from other issue reports. However, these methods showed low performance due to their limited semantic understanding of the issue report text [13].

To overcome this, recent methods have adopted transformer models due to their strong contextual understanding [6, 4, 24, 16, 17]. These studies typically fine-tune the models on historical issues in a supervised setting. For instance, Izadi et al. [16] fine-tuned a RoBERTa [21] model to classify issues into *Bugs*, *Features*, or *Questions*. Similarly, Trautsch et al. [24] trained a seBERT [25] model to classify issues into the same three categories. Although transformer-based approaches achieve strong results, they rely on large amounts of project-specific data to be effective, hindering their generalizability on projects with a small number of labeled issues [13, 16]. Additionally, they struggle to accurately classify minority classes with fewer training examples [17, 6].

These challenges have motivated the exploration of LLMs, which exhibit strong text classification capabilities thanks to extensive pretraining on massive corpora. [13, 3, 2, 5, 26, 8]. Studies show that LLMs perform well in zero-shot settings [8, 2]. SOTA LLM-based methods achieved the best performance by fine-tuning the models on labeled datasets [13, 3]. For example, Aracena et al. [3] instruction-tuned several models on the NLBSE '23 and '24 datasets.

## 3 Approach

### 3.1 Problem Formulation

Given a query issue report, we classify it into only one of the three labels: *bug*, *feature*, or *question*. Only the issue's title and description text are used for classification with no other signals from the issue tracker. The defined label space is adopted directly from prior IRC studies [2, 3, 16, 18, 1].



© Anonymous author(s);

licensed under Creative Commons License CC-BY 4.0

Leibniz International Proceedings in Informatics

LIPICs Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

## 42    3.2    Indexing and Retrieval

43    To prepare the issues for retrieval based on textual similarity, we first embed their textual  
 44    content (i.e., *title + description*). We do not embed the issue labels. Instead, we retain  
 45    them as metadata associated with each embedded issue. To generate the embeddings, we use  
 46    the `all-MiniLM-L6-v2` model from the Sentence-Transformers library [22]. We selected this  
 47    model due to its lightweight architecture, fast performance, and effectiveness on issue report  
 48    text [15]. The embeddings are then indexed using the Faiss vector database library [10],  
 49    which provides efficient high-dimensional similarity retrieval. At the retrieval stage, the query  
 50    issue is first embedded with the same model. Next, Faiss calculates the cosine similarity  
 51    of indexed issue reports to the query, and returns the top-K nearest neighbors, ordered by  
 52    descending similarity.

## 53    3.3    RAGTAG

54    RAGTAG is our proposed retrieval-augmented few-shot prompting approach for IRC. RAG-  
 55    TAG uses the retrieval pipeline described in Section 3.2 to retrieve the top-K nearest neighbors  
 56    to the query issue, and present them as classified examples to the LLM in a few-shot prompt.  
 57    This enables a more accurate classification, based on in-context learning [5]. The choice to  
 58    present the top-K nearest neighbors as examples is based on two insights from prior research:  
 59    First, retrieval-based example selection outperforms random selection [19]. Second, the most  
 60    similar neighbors tend to share the same labels [9]; therefore, they are informative for the  
 61    classification.

62    We create the few-shot prompt in three parts. First, the system message frames the  
 63    classification task and the label space. Second, the top-K nearest neighbors are listed in  
 64    retrieval order (descending similarity), each formatted as *title + description + label*. Finally,  
 65    the query issue is appended to the end of the prompt, ensuring it receives high attention  
 66    from the LLM [20].  $K = 0$  produces a zero-shot prompting baseline, similar to prior work [7].

67    We evaluate RAGTAG at top-K values in  $\{0, 1, 3, 6, 9\}$ . Details for the K value grid  
 68    selection are provided in Section 5.1. Since the median prompt size at  $K = 9$  in our dataset  
 69    is  $\sim 6,300$  tokens, the maximum sequence length is set to 8,192 tokens. Roughly 30% of the  
 70    context quota is reserved for the query issue and 70% for the neighbors. If the neighbors  
 71    exceed their overall quota, they are truncated proportionally based on their length (longer  
 72    neighbors get more truncation). The query issue gets truncated if it exceeds its budget as  
 73    well. Truncation happens on the right side of the description text, preserving all titles and  
 74    labels (if any).

75    The LLM is instructed to generate the predicted label wrapped in XML tags (e.g.  
 76    `<label>bug</label>`). The labels of the presented examples also follow the same format.  
 77    This technique enables easy detection of the predicted label from the LLM’s generation, and  
 78    reduces variance in LLM outputs caused by prompt format sensitivity [23]. We use a regular  
 79    expression to extract the predicted label. Unparseable outputs are marked as *invalid*.

## 80    3.4    VOTAG

81    VOTAG is our proposed non-parametric retrieval IRC method. VOTAG uses the retrieval  
 82    pipeline described in Section 3.2 to retrieve the top-K nearest neighbors of the query issue  
 83    report, then predicts the final label with a similarity-weighted vote on their labels. This  
 84    approach follows the standard K-nearest-neighbor classification setup [9], with neighbors

85 weighted by their similarity to the query [11].<sup>1</sup> Formally, for a query issue  $q$  with retrieved  
 86 neighbors  $\{(n_i, l_i, s_i)\}_{i=1}^K$ , where  $n_i$  is the  $i$ -th nearest neighbor with label  $l_i$  and cosine  
 87 similarity  $s_i$  to  $q$ , the label is predicted by:

$$88 \quad \hat{y} = \arg \max_{c \in \mathcal{L}} \sum_{i=1}^K s_i \cdot \mathbf{1}[l_i = c], \quad (1)$$

89 where  $\mathcal{L}$  is the label set and  $\mathbf{1}[\cdot]$  is the indicator function. VOTAG breaks ties by returning  
 90 the label of the most similar neighbor in the tied classes.

91 VOTAG has three main roles: First, it serves as a fast and cost-efficient baseline that  
 92 LLM-based approaches must clear to justify their computational cost. Second, it is used  
 93 as an ablation to distinguish the contributions of the LLM and retrieval in the RAGTAG  
 94 method. Third, VOTAG’s performance as a function of  $K$  informs the  $K$  value grid evaluated  
 95 for RAGTAG.

### 96 3.5 Fine-Tuning

97 SOTA IRC approaches have achieved strong performance by supervised fine-tuning of  
 98 LLMs [13, 3, 2], using Low-Rank Adaptation (LoRA) [14]. In these studies, the model is  
 99 trained on labeled issue reports and predicts the labels for unseen test issues. Each training  
 100 example is formatted as an instruction prompt which includes the issue’s title and description  
 101 as input, and its label as output. At inference, the trained model receives the test issue in  
 102 the same format, and predicts the empty label field. Regular expressions are used to extract  
 103 the prediction from the output, and unparseable outputs are marked as *invalid*. We use this  
 104 established methodology as our fine-tuning baseline.

105 For memory efficiency, each model is loaded in 4-bit quantization with Unsloth<sup>2</sup>. The  
 106 LoRA rank and alpha are set to 16, with a learning rate of  $2 \times 10^{-4}$  and the paged AdamW  
 107 8-bit optimizer. Training runs for 3 epochs with a per-device batch size of 1 and gradient  
 108 accumulation steps of 16. The max sequence length is set to 2,048 tokens, which covers  
 109 94.9% of the issues in our dataset. Issues that exceed this budget get truncated from the  
 110 right side of their description text to fit.

## 111 4 Experimental Setup

### 112 4.1 Dataset and Settings

113 We use the dataset introduced in Heo et al. [13] for our experiments. The dataset has a total  
 114 of 6,600 issue reports gathered from eleven active open-source projects: `ansible/ansible`,  
 115 `bitcoin/bitcoin`, `dart-lang/sdk`, `dotnet/roslyn`, `facebook/react`, `flutter/flutter`,  
 116 `kubernetes/kubernetes`, `microsoft/TypeScript`, `microsoft/vscode`, `opencv/opencv`, and  
 117 `tensorflow/tensorflow`. Each project has 600 issue reports with an equal distribution  
 118 across the labels (*bug*, *feature*, *question*). To support their LLM fine-tuning method, Heo et  
 119 al. divided the dataset into 50/50 train and test splits, stratified by project and label. This  
 120 results in 300 train and 300 test issue reports per project.

121 We also adopt Heo et al.’s project-specific (PS) and project-agnostic (PA) configurations  
 122 to examine how data scope affects IRC performance. In the PS setting, each project’s train

<sup>1</sup> Squared-similarity weighting and uniform majority were also evaluated, but were outperformed by similarity-weighted voting.

<sup>2</sup> <https://unsloth.ai>

and test data are used in isolation, resulting in eleven different train-test pairs. In the PA setting, the train splits of all projects are merged into a single training set of 3,300 issue reports. In both configurations, RAGTAG and VOTAG only use the training splits as their retrieval index to ensure fair comparison with the fine-tuning baseline.

## 4.2 Models

We use the Qwen2.5-Instruct model family [27] in our evaluations with four different parameter sizes: 3B, 7B, 14B, and 32B. Qwen2.5 is an open-source LLM with strong performance on general-purpose benchmarks. Using the same model at different parameter sizes isolates the effect of model scale from model architecture across LLM-based IRC approaches. Model temperature is set at 0.1 for deterministic classifications.

## 4.3 Evaluation Metrics

We report macro-averaged precision, recall, and  $F_1$  over the three labels (**bug**, **feature**, **question**), along with overall accuracy. Per-class metrics are reported when the asymmetry across classes is itself the subject of analysis. Outputs that fail to parse a valid label (after the XML-tag extraction and regex fallback described in ??) are counted as incorrect, and an *invalid rate* is reported separately so any parsing-induced loss is auditable.

### 4.3.0.1 Aggregation convention (PA vs PS).


The PA setting produces a single set of predictions over the full 3,300-issue test set, and we compute macro  $F_1$  directly over that set. The PS setting produces eleven sets of predictions, one per project, each over its own  $\sim 300$ -issue test slice. Throughout the paper we report PS results using *pooled* aggregation: the eleven per-project prediction sets are concatenated into one 3,300-issue test set and macro  $F_1$  is computed once on the pool. This treats each test issue as one observation in both settings and yields apples-to-apples comparisons between PA and PS. The alternative — computing macro  $F_1$  separately within each project and averaging the eleven numbers — weights projects rather than issues; because macro  $F_1$  is non-linear in the underlying confusion-matrix counts, the two conventions disagree, and on our data the per-project mean is consistently 0.002–0.011 macro  $F_1$  below the pooled value with the gap largest for fine-tuning. We adopt pooling as the canonical reporting convention to remove this asymmetry; per-project breakdowns appear only where the project is the unit of analysis and are explicitly labelled as such.

## 4.4 Setup Hardware

# 5 Evaluations

## 5.1 RQ1) To what extent can similarity-weighted voting on the nearest neighbors (VOTAG) classify issue reports, and at what point does adding neighbors stop helping?

We evaluate VOTAG’s performance on both PA and PS settings at  $k$  values 1 to 30. The resulting macro F1 scores are shown in Figure 1. VOTAG’s best macro F1 is 0.595 ( $k=15$ ) in PS and 0.604 ( $k=16$ ) in PA. Both settings reach 99.8% (PS) and 98.5% (PA) of their respective peak by  $k=7$ . When adding neighbors beyond the peak, PS loses 0.015 and PA loses 0.011 by  $k=30$ .



■ **Figure 1 (left)** VOTAG macro  $F_1$  as a function of  $k$  in both PS and PA settings. Circles mark peak performance. **(right)** Per-class  $F_1$  at each setting’s best  $k$  ( $k=15$  for PS,  $k=16$  for PA).

PA consistently outperforms PS at every  $k$ . However, the performance gap is marginal: averaging 0.007 macro  $F_1$ , with a maximum of 0.015 at  $k=18$ . This similar performance is due to the large overlap of the top- $k$  neighbors in both settings. Specifically, PA and PS share 84–90% of the top- $k$  neighbors across  $k \in [3, 30]$ . Therefore, even when using the full 3,300-issue retrieval index, most of the retrieved neighbors are from the same project. This means issue reports from the same project tend to be closer in the embedding space.

Question consistently has the weakest performance among the three labels at every  $k$  in both settings. At best  $k$ , PS per-class  $F_1$  is 0.630 for bug, 0.614 for feature, and 0.541 for question. The corresponding PA values are 0.632, 0.629, and 0.550 (Figure 1). Additionally, questions are more frequently misclassified as bugs: 32% of question issues are predicted as bugs in both settings, while 18% (PS) and 17% (PA) are predicted as features. Since VOTAG predicts the label with the most accumulated similarity weight, this misclassification asymmetry indicates that question issues are more similar to bugs than to features in the embedding space.


## 5.2 RAGTAG

We evaluate RAGTAG at  $k$  values 0, 1, 3, 6, 9, 12, 15 on both PA and PS settings, where  $k=0$  is the zero-shot baseline [7].  $k$  value grid selection is informed by VOTAG’s peak results around  $k=15$  in the previous section. Figure 2 shows the macro  $F_1$  of every model across this grid for both settings.

PS marginally outperforms PA at almost every (model,  $k$ ) pair evaluated. The close performance gap is due to the same top- $k$  neighbor overlap observed in the VOTAG analysis. However, This marginal improvement is significant because PS uses  $11 \times$  less retrieval data per project: each PS index covers only its own 300 retrieval issues, while PA uses the full 3,300 across all 11 projects. Given this practical advantage, we conduct the rest of our RAGTAG analysis on the PS setting and report PS numbers throughout this section.

The best  $k$  value grows nearly monotonically with model size. Qwen-3B peaks at  $k=3$  (macro  $F_1 = 0.697$ ), Qwen-7B at  $k=6$  (0.718), and both Qwen-14B and Qwen-32B at  $k=12$  (0.732 and 0.767 respectively). This shows that larger models can use the context from additional examples more effectively. However, adding neighbors past  $k=12$  has no benefit and slightly degrades performance. Every model loses 0.001–0.015 macro  $F_1$  going from  $k=12$  to  $k=15$ .

Every RAGTAG configuration outperforms VOTAG’s best (0.604 at PA  $k=16$ ), with



**Figure 2** RAGTAG macro  $F_1$  as a function of  $k$  for all four Qwen sizes. Solid lines with filled markers are project-specific (PS); dashed lines with hollow markers are project-agnostic (PA).  $k=0$  is the zero-shot baseline. The dotted horizontal line marks VOTAG’s best (0.604, PA  $k=16$ ).

the smallest margin from Qwen-3B zero-shot ( $F_1 = 0.613$ ,  $+0.009$ ) and the largest from Qwen-32B at  $k=12$  ( $F_1 = 0.767$ ,  $+0.163$ ). Overall, this shows that adding the LLM’s semantic understanding can meaningfully improve IRC on top of pure retrieval. However, It is impressive that VOTAG, a cost-efficient retrieval-only heuristic with no LLM reasoning, comes within 0.009 macro  $F_1$  of Qwen-3B zero-shot.

## 6 Discussion

## 7 Threats to Validity

## 8 Conclusion

## 9 Data Availability

## References

- Giuliano Antoniol, Kamel Ayari, Massimiliano Di Penta, Foutse Khomh, and Yann-Gaël Guéhéneuc. Is it a bug or an enhancement? a text-based approach to classify change requests. In *Proceedings of the 2008 conference of the center for advanced studies on collaborative research: meeting of minds*, pages 304–318, 2008.
- Gabriel Aracena, Kyle Luster, Fabio Santos, Igor Steinmacher, and Marco A. Gerosa. Applying large language models api to issue classification problem, 2024. URL: <https://arxiv.org/abs/2401.04637>, arXiv:2401.04637.
- Gabriel Aracena, Kyle Luster, Fabio Santos, Igor Steinmacher, and Marco A Gerosa. Applying large language models to issue classification: Revisiting with extended data and new models. *Science of Computer Programming*, page 103333, 2025.
- Shikhar Bharadwaj and Tushar Kadam. Github issue classification using bert-style models. In *Proceedings of the 1st International Workshop on Natural Language-based Software Engineering*, pages 40–43, 2022.

- 218   **5**   Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla  
219       Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language  
220       models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901,  
221       2020.
- 222   **6**   Giuseppe Colavito, Filippo Lanubile, and Nicole Novielli. Issue report classification using  
223       pre-trained language models. In *Proceedings of the 1st International Workshop on Natural  
224       Language-based Software Engineering*, pages 29–32, 2022.
- 225   **7**   Giuseppe Colavito, Filippo Lanubile, and Nicole Novielli. Benchmarking large language models  
226       for automated labeling: The case of issue report classification. *Information and Software  
227       Technology*, page 107758, 2025.
- 228   **8**   Giuseppe Colavito, Filippo Lanubile, Nicole Novielli, and Luigi Quaranta. Leveraging gpt-like  
229       llms to automate issue labeling. In *Proceedings of the 21st International Conference on Mining  
230       Software Repositories*, pages 469–480, 2024.
- 231   **9**   Thomas Cover and Peter Hart. Nearest neighbor pattern classification. *IEEE transactions on  
232       information theory*, 13(1):21–27, 1967.
- 233   **10**   Matthijs Douze, Alexandr Guzhva, Chengqi Deng, Jeff Johnson, Gergely Szilvasy, Pierre-  
234       Emmanuel Mazaré, Maria Lomeli, Lucas Hosseini, and Hervé Jégou. The faiss library. *IEEE  
235       Transactions on Big Data*, 2025.
- 236   **11**   Sahibsingh A Dudani. The distance-weighted k-nearest-neighbor rule. *IEEE Transactions on  
237       Systems, Man, and Cybernetics*, (4):325–327, 1976.
- 238   **12**   Qiang Fan, Yue Yu, Gang Yin, Tao Wang, and Huaimin Wang. Where is the road for issue  
239       reports classification based on text mining? In *2017 ACM/IEEE International Symposium on  
240       Empirical Software Engineering and Measurement (ESEM)*, pages 121–130. IEEE, 2017.
- 241   **13**   Jueun Heo and Seonah Lee. A study on applying large language models to issue classification.  
242       In *2025 IEEE/ACM 33rd International Conference on Program Comprehension (ICPC)*, pages  
243       1–11. IEEE Computer Society, 2025.
- 244   **14**   Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Liang  
245       Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *Iclr*, 1(2):3,  
246       2022.
- 247   **15**   Huihui Huang, Ratnadira Widyasari, Ting Zhang, Ivana Clairine Irsan, Jieke Shi, Han Wei  
248       Ang, Frank Liauw, Eng Lieh Ouh, Lwin Khin Shar, Hong Jin Kang, et al. Back to the basics:  
249       Rethinking issue-commit linking with llm-assisted retrieval. *arXiv preprint arXiv:2507.09199*,  
250       2025.
- 251   **16**   Maliheh Izadi. Catiss: An intelligent tool for categorizing issues reports using transformers. In  
252       *Proceedings of the 1st international workshop on natural language-based software engineering*,  
253       pages 44–47, 2022.
- 254   **17**   Maliheh Izadi, Kiana Akbari, and Abbas Heydarnoori. Predicting the objective and priority  
255       of issue reports in software repositories. *Empirical Software Engineering*, 27(2):50, 2022.
- 256   **18**   Rafael Kallis, Andrea Di Sorbo, Gerardo Canfora, and Sebastiano Panichella. Ticket tagger:  
257       Machine learning driven issue classification. In *2019 IEEE International Conference on  
258       Software Maintenance and Evolution (ICSME)*, pages 406–409. IEEE, 2019.
- 259   **19**   Jiachang Liu, Dinghan Shen, Yizhe Zhang, William B Dolan, Lawrence Carin, and Weizhu  
260       Chen. What makes good in-context examples for gpt-3? In *Proceedings of Deep Learning  
261       Inside Out (DeeLIO 2022): The 3rd workshop on knowledge extraction and integration for  
262       deep learning architectures*, pages 100–114, 2022.
- 263   **20**   Nelson F Liu, Kevin Lin, John Hewitt, Ashwin Paranajpe, Michele Bevilacqua, Fabio Petroni,  
264       and Percy Liang. Lost in the middle: How language models use long contexts. *Transactions  
265       of the association for computational linguistics*, 12:157–173, 2024.
- 266   **21**   Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy,  
267       Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized bert  
268       pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019.

- 269    **22**    Nils Reimers and Iryna Gurevych. Sentence-bert: Sentence embeddings using siamese bert-  
270           networks. *arXiv preprint arXiv:1908.10084*, 2019.
- 271    **23**    Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. Quantifying language models’  
272           sensitivity to spurious features in prompt design or: How i learned to start worrying about  
273           prompt formatting. *arXiv preprint arXiv:2310.11324*, 2023.
- 274    **24**    Alexander Trautsch and Steffen Herbold. Predicting issue types with sebert. In *Proceedings*  
275           *of the 1st international workshop on natural language-based software engineering*, pages 37–39,  
276           2022.
- 277    **25**    Julian Von der Mosel, Alexander Trautsch, and Steffen Herbold. On the validity of pre-trained  
278           transformers for natural language processing in the software engineering domain. *IEEE*  
279           *Transactions on Software Engineering*, 49(4):1487–1507, 2022.
- 280    **26**    Jason Wei, Maarten Bosma, Vincent Y Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan  
281           Du, Andrew M Dai, and Quoc V Le. Finetuned language models are zero-shot learners. *arXiv*  
282           *preprint arXiv:2109.01652*, 2021.
- 283    **27**    An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan  
284           Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2.5 technical report. *arXiv preprint*  
285           *arXiv:2412.15115*, 2024.